

NLP Techniques and Applications

- Text prediction using RNN

N-Grams

- Large N-grams to capture dependencies between distant words
- Need a lot of space and RAM
 - N-grams consume a lot of memory
- Different types of RNNs are the preferred alternative

Recurrent Neural Networks

Nour was supposed to study with me. I called her but she did not answer
have


-----want-----
respond
choose
want
have
ask
attempt
answer
know

RNNs look at every previous word

Similar probabilities with trigram


Recurrent Neural Networks

I called her but she did not _____




Recurrent Neural Networks

I called her but she did not _____



Recurrent Neural Networks




Q. An RNN would have the same number of parameters for word sequences of different lengths?

Recurrent Neural Networks

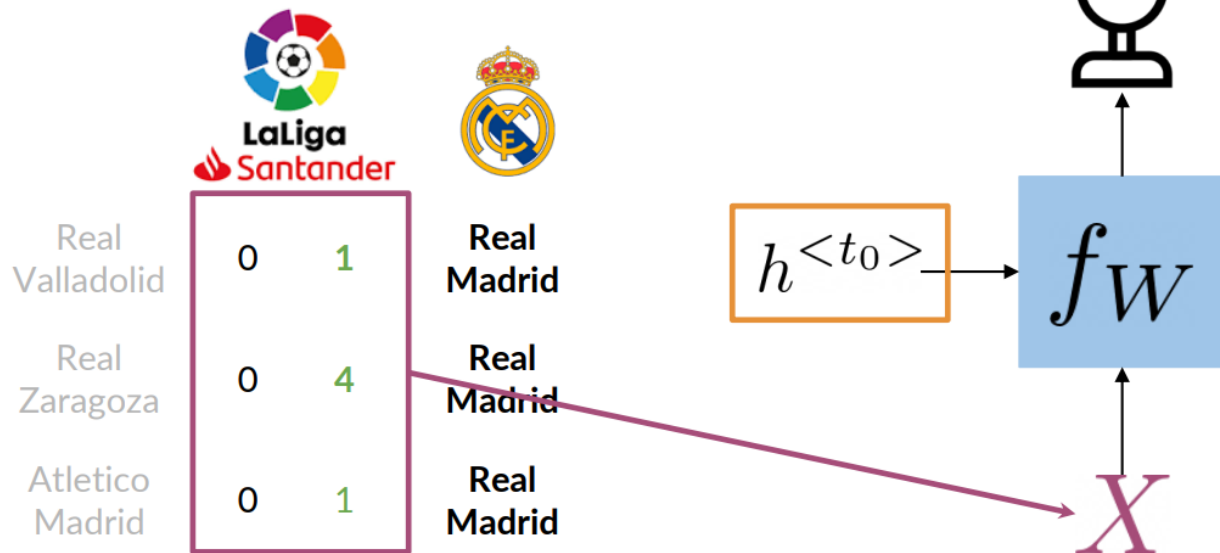
Summary

- RNNs model relationships among distant words
- In RNNs a lot of computations share parameters



Applications of RNN


One to One



Applications of RNN

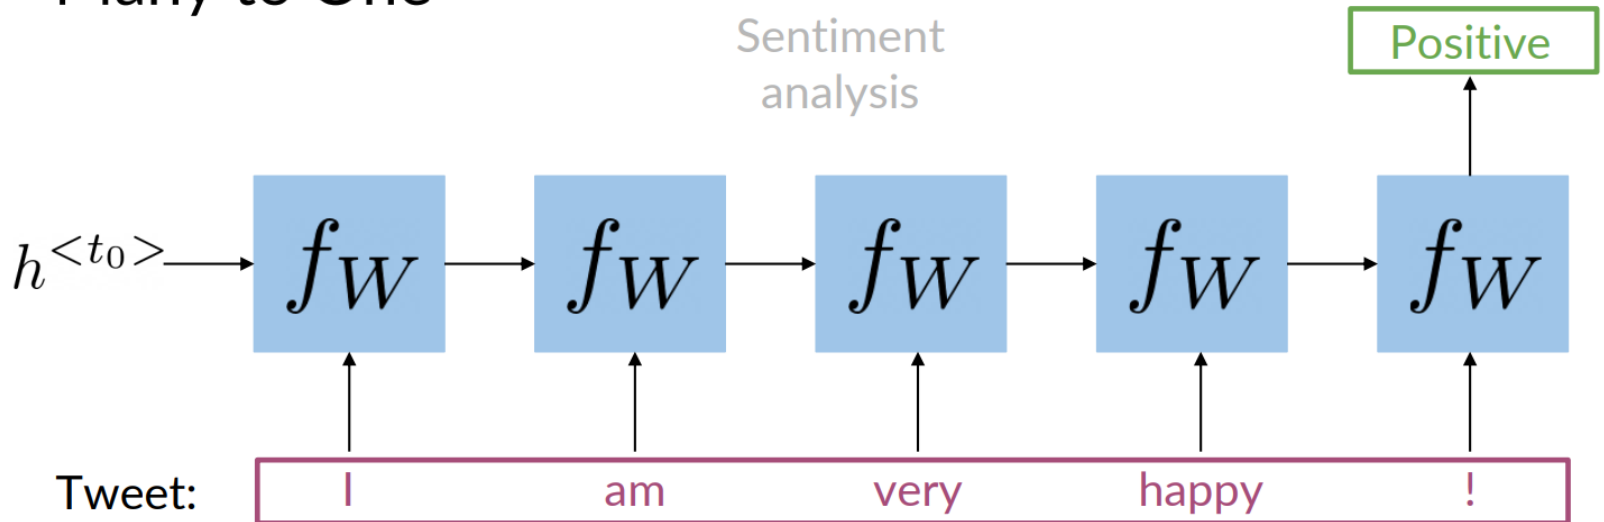
One to Many

Caption
generation



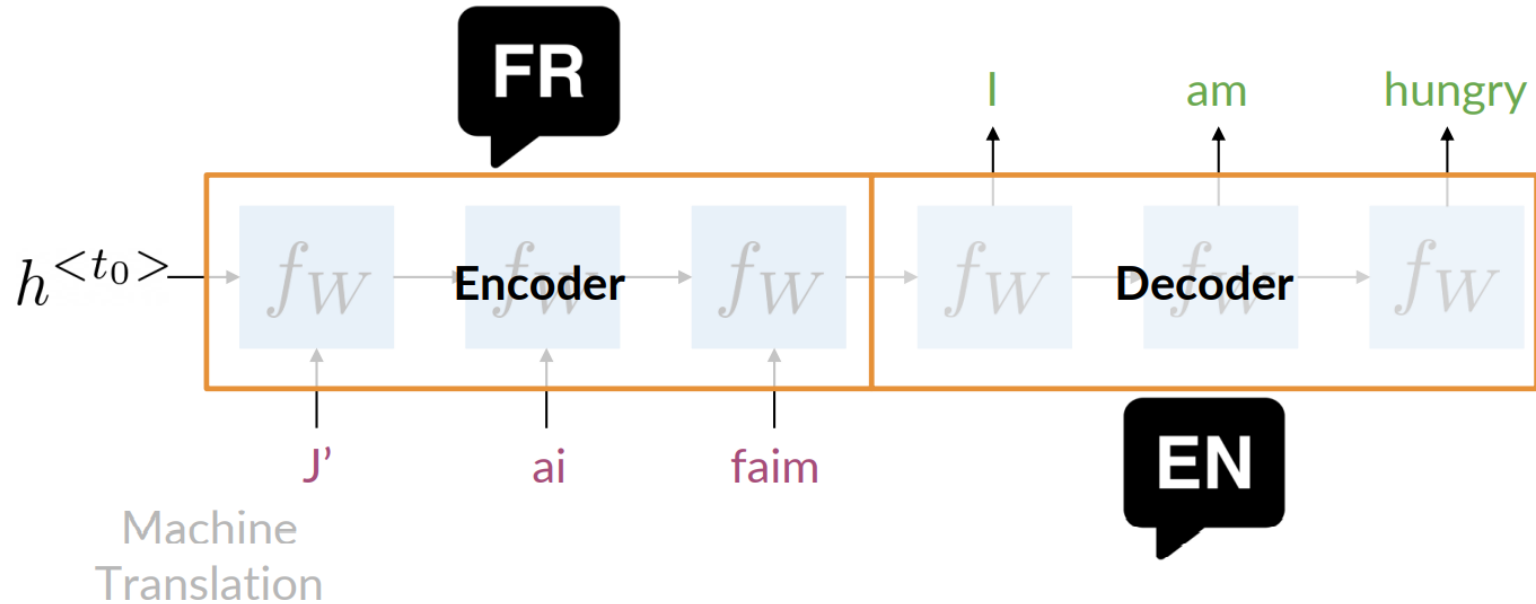
Applications of RNN

Many to One



Applications of RNN

Many to Many



Applications of RNN

Q. “one to many” architecture?

- An RNN which inputs a conversation and determines the topic.
- An RNN which inputs a topic and generates a conversation about that topic.
- An RNN which inputs a sentiment and generates a sentence.
- An RNN which inputs a sentence and determines the sentiment.

RNN

Summary


- RNNs can be implemented for a variety of NLP tasks
- Applications include Machine translation and caption generation

Maths in simple RNN

- How RNNs propagate information (Through time!)
- How RNNs make predictions

Maths in simple RNN



A Vanilla RNN



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$
$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$



$$h^{<t>} = g(W_{hh}h^{<t-1>} \oplus W_{hx}x^{<t>} + b_h)$$

Maths in simple RNN




$$h^{<t>} = g(W_{hh}h^{<t-1>} + W_{hx}x^{<t>} + b_h)$$

Maths in simple RNN




$$h^{<t>} = g(\boxed{W_{hh}h^{<t-1>} + W_{hx}x^{<t>}} + b_h)$$

Maths in simple RNN



$$h^{<t>} = g(W_{hh}h^{<t-1>} + W_{hx}x^{<t>} + b_h)$$

Maths in simple RNN



Note that we end up training with : $W_{hh}, W_{hx}, W_{yh}, b_h, b_y$


Maths in simple RNN

Summary:

- Hidden states propagate information through time
- Basic recurrent units have two inputs at each time: $h^{<t-1>}$ $x^{<t>}$

Cost Function for RNNs

Cross Entropy Loss



K - classes or possibilities


$$J = - \sum_{j=1}^K y_j \log \hat{y}_j$$

Either 0 or 1

Looking at a single example (x, y)

Cost Function for RNNs

Cross Entropy Loss



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$

$$J = -\frac{1}{T} \sum_{t=1}^T \sum_{j=1}^K y_j^{<t>} \log \hat{y}_j^{<t>}$$

Average with respect to time

For RNNs the loss function is just an average through time!